

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.linear_model import LinearRegression
from sklearn.pipeline import make_pipeline
from sklearn.metrics import r2_score

import statsmodels.api as sm
import statsmodels.formula.api as smf

import scipy.stats as stats
from sklearn.preprocessing import RobustScaler
import warnings
warnings.filterwarnings("ignore")

drop_cols = [
    "TWD", "COG", "TWA", "TimeUTC", "SecondsSince1970", "ISODateTimeUTC", "Latitude", "Longitude", "Euler_X (deg)", "Euler_Y (deg)",
    "Lat", "LatBow", "LatCenter", "LatStern", "Lon", "LonBow", "LonCenter", "LonStern",
    "Leg", "Log", "LogAlongCourse", "MagneticVariation", "Rank", "TimeLocal",
    "DistanceToLeader", "interval_id", "boat_name", "interval_duration",
    "Heel", "Heel_Lwd", "Line_R", "Line_L", "BelowLineCalc", "VMC", "XTE", "VMG", "gain_forward", "gain_lateral", "gain_vmg", "Total_lines", "LoadCell_1", "LoadCell_2", "LoadCell_3", "LoadCell_4"
]

def show_correlation_matrix(dataframe, taille_figure=(6, 4), cmap="coolwarm"):
    plt.figure(figsize=taille_figure)
    sns.heatmap(dataframe.corr(), cmap=cmap, center=0)
    plt.title("Correlation matrix")
    plt.show()

def show_target_correlation(df, variable="SOG", taille_figure=(6, 4)):
    corr_with_vmg = df.corr()[variable].sort_values(ascending=False)
    #print(f"Correlation with {variable} :")
    # print(corr_with_vmg)

    plt.figure(figsize=taille_figure)
    sns.heatmap(corr_with_vmg.to_frame(), annot=True, cmap="coolwarm")
    plt.title(f"Correlation with {variable}")
    plt.show()

def compute_anova(df, target="SOG"):
    # Make a copy of the DataFrame to avoid modifying the original
    df_copy = df.copy()
    cols = df_copy.columns.drop(target)
    df_no_missing = df_copy.dropna(subset=cols)
    if df_no_missing.empty:
        print("⚠ Pas assez de données non-nulles pour faire une ANOVA.")
        return pd.DataFrame() # retourne un DataFrame vide ou None
    formula_no_missing = f"{target} ~ " + " + ".join(
        list(df_no_missing.columns.drop(target)))
    model_no_missing = smf.ols(formula=formula_no_missing, data=df_no_missing).fit()
    anova_no_missing = sm.stats.anova_lm(model_no_missing, typ=2)
```

```

        return anova_no_missing.sort_values("F", ascending=False)

# With simple standard scaler
def linear_regression(df, target="SOG", degree=1, top_coefs=30):
    df_copy = df.copy()
    cols = df_copy.columns.drop(target)

    df_clean = df_copy.dropna(subset=cols)
    if df_clean[target].isna().any():
        df_clean = df_clean.dropna(subset=[target])

    X = df_clean[cols]
    y = df_clean[target]

    poly = PolynomialFeatures(degree=degree, include_bias=False)
    X_poly = poly.fit_transform(X)
    feature_names = poly.get_feature_names_out(X.columns)

    model = make_pipeline(
        StandardScaler(),
        LinearRegression()
    )

    model.fit(X_poly, y)
    y_pred = model.predict(X_poly)

    # Affichage de la performance
    print(f"R²: {r2_score(y, y_pred):.3f}")
    print(f"Used samples: {len(X)}")

    # Coefficients
    coefs = model.named_steps['linearregression'].coef_
    intercept = model.named_steps['linearregression'].intercept_

    coef_df = pd.DataFrame({
        'feature': feature_names,
        'coefficient': coefs
    }).sort_values('coefficient', key=abs, ascending=False)

    # Formule affichée
    print("\n Formula :")
    terms = [f"{coef:.3f} * {feat}" for feat, coef in zip(feature_names, coefs)]
    equation = " + ".join(terms)
    print(f"\n {target} ≈ {intercept:.3f} + {equation}\n")

    return coef_df.head(top_coefs)
"""

from sklearn.linear_model import RidgeCV
# Regularisaton with lasso
def linear_regression(df, target="SOG", top_coefs=30):
    df_copy = df.copy()
    cols = df_copy.columns.drop(target)

    # Drop rows with missing values
    df_clean = df_copy.dropna(subset=cols)
    if df_clean[target].isna().any():
        df_clean = df_clean.dropna(subset=[target])

```

```

X = df_clean[cols]
y = df_clean[target]

# Ridge regression with standard scaling
model = make_pipeline(
    StandardScaler(),
    RidgeCV(alphas=[0.1, 1.0, 10.0])
)

model.fit(X, y)
y_pred = model.predict(X)

# Performance
print(f"R²: {r2_score(y, y_pred):.3f}")
print(f"Used samples: {len(X)}")

# Extract coefficients
coefs = model.named_steps['ridgecv'].coef_
intercept = model.named_steps['ridgecv'].intercept_

coef_df = pd.DataFrame({
    'feature': X.columns,
    'coefficient': coefs
}).sort_values('coefficient', key=abs, ascending=False)

# Display formula
print("\nFormula:")
terms = [f"{coef:.3f} * {feat}" for feat, coef in zip(X.columns, coefs)]
equation = " + ".join(terms)
print(f"\n{target} ~ {intercept:.3f} + {equation}\n")

return coef_df.head(top_coefs)

from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import RobustScaler
from sklearn.metrics import r2_score

def linear_regression(df, target="SOG", top_coefs=30):
    df_copy = df.copy()

    # Drop missing values
    df_clean = df_copy.dropna(subset=[target] + [col for col in df_copy.columns if col != target])
    X = df_clean.drop(columns=[target])
    y = df_clean[target]

    # Manual robust scaling
    scaler = RobustScaler()
    X_scaled = pd.DataFrame(scaler.fit_transform(X), columns=X.columns)

    # OLS Regression
    model = LinearRegression()
    model.fit(X_scaled, y)
    y_pred = model.predict(X_scaled)

    # Model performance
    print(f"R²: {r2_score(y, y_pred):.3f}")
    print(f"Used samples: {len(X_scaled)}")

```

```

# Coefficient analysis
coefs = model.coef_
intercept = model.intercept_

coef_df = pd.DataFrame({
    "feature": X.columns,
    "coefficient": coefs
}).sort_values("coefficient", key=abs, ascending=False)

# Print readable formula
print("\nFormula:")
terms = [f"{coef:.3f} * {feat}" for feat, coef in zip(X.columns, coefs)]
equation = " +\n ".join(terms)
print(f"\n{target} ≈ {intercept:.3f} +\n {equation}\n")

return coef_df.head(top_coefs)
"""
"""

def t_test(df1, df2, target="SOG"):
    t_stat, p_value = stats.ttest_ind(df1[target].dropna(), df2[target].dropna())
    print(f"\nT-statistic: {t_stat:.3f}, p-value: {p_value:.15f}\n")

    # If p-value is less than 0.05, the difference is statistically significant
    if p_value < 0.05:
        print("The difference is statistically significant, keeping data split.")
    else:
        print("The difference is not statistically significant, keeping data combined.\n")

def full_analysis(df_numeric, target_variable="SOG"):
    # Drop 'Line_Side' column only if it's fully NaN
    if "Side_lines" in df_numeric.columns and df_numeric["Side_lines"].isna().all():
        df_numeric = df_numeric.drop(columns=["Side_lines"])
        print("Dropped column 'Side_lines' (all values were NaN).")

    # Display correlation with the target variable
    print(f"\nCorrelation with {target_variable}:")
    show_target_correlation(df_numeric, variable=target_variable)

    # Compute and display ANOVA results
    print("\nANOVA:")
    anova_results = compute_anova(df_numeric)
    display(anova_results)

    # Apply and display polynomial regression results
    print("\nPolynomial fit:")
    regression_results = linear_regression(df_numeric)
    display(regression_results)

```

```

In [2]: df = pd.read_csv("all_data_enriched.csv")
df = df[df["boat_name"] == "SenseBoard"].copy()
load_cell_cols = ["LoadCell_1", "LoadCell_2", "LoadCell_3", "LoadCell_4", "LoadCell_5", "LoadCell_6"]
df = df.dropna(subset=load_cell_cols, how='all')

df.sample(10)

```

Out[2]:

| | | ISODateTimeUTC | SecondsSince1970 | Heel_Abs | Heel_Lwd | Lat | LatBow | LatCenter | LatStern | Leg | Line_C | ... | M_tot_X | M_tot_Y | M_front_X | M_front_Y | M_back_X | M_back_Y |
|-------|--|-------------------------------------|------------------|----------|----------|-----------|-----------|-----------|-----------|-----|---------|-----|--------------|-------------|--------------|-------------|--------------|----------|
| 27182 | | 2025-06-07 12:18:12.359000+00:00 | 1.749299e+09 | 38.3 | 38.3 | 43.514855 | 43.514857 | 43.514851 | 43.514846 | 1.0 | 67.000 | ... | -1283.160374 | -284.229902 | -815.910320 | -92.576328 | -467.250054 | 30 |
| 55595 | | 2025-06-09 13:28:07.360000+00:00 | 1.749476e+09 | 47.6 | 47.6 | 43.506216 | 43.506218 | 43.506212 | 43.506206 | NaN | 5.300 | ... | -2371.706268 | 1418.689814 | -1040.980773 | -197.525690 | -1330.725495 | -89 |
| 63645 | | 2025-06-10 12:44:24.857000+00:00 | 1.749559e+09 | 53.3 | 53.3 | 43.534359 | 43.534357 | 43.534363 | 43.534369 | NaN | 115.715 | ... | 10011.516880 | 1616.431532 | 7203.701558 | 1764.848024 | 2807.815321 | -36 |
| 68295 | | 2025-06-10 13:02:46.457000+00:00 | 1.749561e+09 | 56.3 | 56.3 | 43.533271 | 43.533269 | 43.533275 | 43.533281 | NaN | 101.400 | ... | 10863.739839 | 7005.677461 | 6543.846389 | 1262.898776 | 4319.893450 | -134 |
| 34502 | | 2025-06-07 13:03:47.063000+00:00 | 1.749301e+09 | 50.4 | 50.4 | 43.514103 | 43.514105 | 43.514100 | 43.514096 | 1.0 | 77.076 | ... | -777.176717 | -765.271102 | -864.138144 | -122.082255 | 86.961427 | -8 |
| 34468 | | 2025-06-07 13:03:45.357000+00:00 | 1.749301e+09 | 33.7 | 33.7 | 43.513952 | 43.513953 | 43.513948 | 43.513944 | 1.0 | 78.200 | ... | -1272.297894 | -255.606557 | -1437.347670 | -123.292255 | 165.049777 | -10 |
| 78715 | | 2025-06-10 13:43:12.553000+00:00 | 1.749563e+09 | 45.7 | 45.7 | 43.535173 | 43.535175 | 43.535169 | 43.535163 | NaN | 93.900 | ... | -1079.209155 | 4447.514202 | -508.607748 | 296.444850 | -570.601407 | -97 |
| 26537 | | 2025-06-07 12:15:31.444000+00:00 | 1.749299e+09 | 58.6 | 58.6 | 43.516040 | 43.516038 | 43.516044 | 43.516049 | 1.0 | 117.100 | ... | 5743.769612 | -393.365826 | 4733.782084 | 743.396538 | 1009.987527 | -33 |
| 66590 | | 2025-06-10 12:53:38.471000+00:00 | 1.749560e+09 | 65.6 | 65.6 | 43.531291 | 43.531289 | 43.531295 | 43.531300 | NaN | 109.200 | ... | 7106.104112 | 7502.418021 | 5057.937720 | 334.410612 | 2048.166392 | -62 |
| 72213 | | 2025-06-10 13:22:39.556000+00:00 | 1.749562e+09 | 56.9 | 56.9 | 43.534803 | 43.534801 | 43.534807 | 43.534812 | NaN | 144.000 | ... | 11601.677276 | 6143.395780 | 7424.225717 | 1738.473400 | 4177.451558 | -134 |

10 rows × 71 columns



I. All together

```
# Select numeric columns
df_numeric = df.select_dtypes(include=["float64", "int64"]).copy()

# Drop specified columns
df_numeric.drop(columns=[c for c in drop_cols if c in df_numeric.columns], inplace=True)

# Drop rows where 'SOG' is missing
df_numeric.dropna(subset=["SOG"], inplace=True)

# Print summary
print(f"Variables utilisées: {df_numeric.columns.tolist()}")
print(f"Number of rows after filtering: {len(df_numeric)}")
```

Variables utilisées: ['Heel_Abs', 'Line_C', 'ROT', 'Side_lines', 'Trim', 'TWA_Abs', 'TWS', 'SOG', 'boat_weight', 'F_front', 'F_back']
Number of rows after filtering: 30546

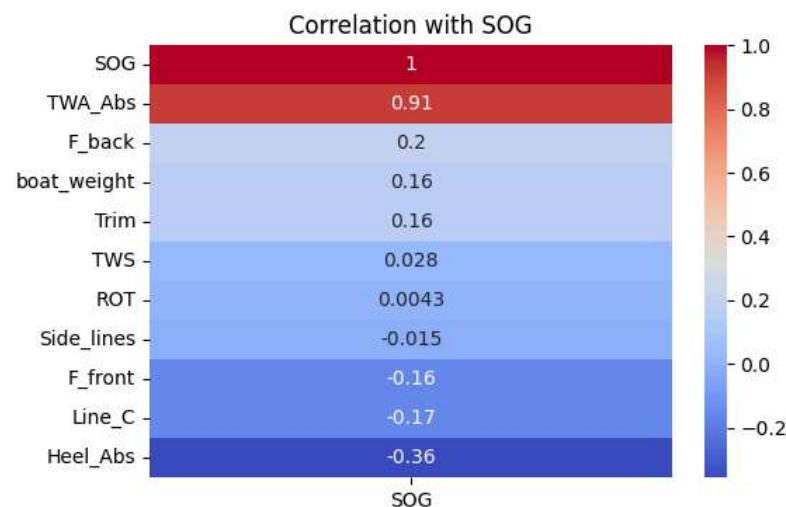
In [4]: df_numeric.sample(20)

Out[4]:

| | Heel_Abs | Line_C | ROT | Side_lines | Trim | TWA_Abs | TWS | SOG | boat_weight | F_front | F_back |
|--------------|----------|---------|---------|------------|------|---------|--------|------|-------------|--------------|--------------|
| 62817 | 53.2 | 90.100 | 4.587 | NaN | 7.8 | 130.200 | 6.866 | 24.6 | 102.89 | 34784.008390 | 53934.845417 |
| 62370 | 59.1 | 83.500 | 4.545 | NaN | 14.1 | 126.665 | 7.165 | 25.4 | 102.89 | 37443.014849 | 56549.027043 |
| 19652 | 55.4 | 115.092 | -36.893 | 4.500 | 6.3 | 52.200 | 11.856 | 23.6 | 109.09 | 11881.282987 | 27775.947539 |
| 32198 | 49.1 | 7.900 | -9.783 | NaN | 9.6 | 144.107 | 7.758 | 26.0 | 102.89 | 28630.250976 | 34016.103360 |
| 34549 | 48.1 | 87.358 | -3.960 | 11.600 | 8.0 | 143.617 | 8.800 | 26.6 | 109.09 | 23115.218456 | 27020.527284 |
| 19480 | 46.1 | 100.200 | -38.000 | 6.827 | 7.6 | 52.084 | 11.200 | 22.8 | 109.09 | 13793.963897 | 32081.398393 |
| 20552 | 44.3 | 89.200 | 6.000 | 15.000 | 9.6 | 157.000 | 9.411 | 28.6 | 109.09 | 19012.873518 | 41587.145802 |
| 50972 | 49.3 | 2.600 | -7.292 | 92.863 | 6.6 | 141.296 | 9.465 | 27.8 | 115.09 | 24486.993259 | 42587.045290 |
| 68552 | 59.9 | 119.192 | 2.174 | NaN | 6.4 | 59.000 | 5.600 | 22.7 | 102.89 | 43521.757989 | 58379.976922 |
| 56028 | 53.1 | 4.500 | 24.490 | 130.200 | 11.0 | 47.048 | 7.774 | 22.3 | 115.09 | 34156.267070 | 43871.435195 |
| 56202 | 58.1 | 6.600 | -7.071 | 116.700 | 6.8 | 51.691 | 7.500 | 21.8 | 115.09 | 36110.393162 | 40964.210846 |
| 28885 | 58.7 | 5.300 | 19.608 | NaN | 8.2 | 28.726 | 8.600 | 20.5 | 102.89 | 27353.351567 | 20124.249880 |
| 63237 | 64.7 | 109.600 | 3.738 | NaN | 10.4 | 51.600 | 6.000 | 21.5 | 102.89 | 37185.222885 | 47818.874979 |
| 49321 | 54.7 | 9.500 | 20.952 | 130.341 | 9.3 | 52.816 | 6.400 | 23.8 | 115.09 | 39280.990480 | 35715.303430 |
| 54763 | 37.3 | 5.300 | 17.431 | 92.330 | 4.7 | 58.736 | 6.864 | 23.6 | 115.09 | 34150.443219 | 34482.984248 |
| 12027 | 57.1 | 128.822 | 7.865 | 15.029 | 8.7 | 42.025 | 9.000 | 21.7 | 109.09 | 20042.851331 | 33621.791142 |
| 37906 | 62.9 | 117.200 | -9.184 | 13.622 | 9.3 | 35.556 | 10.955 | 23.2 | 109.09 | 23649.054986 | 28524.625068 |
| 62934 | 54.7 | 127.700 | -23.913 | NaN | 4.6 | 129.300 | 6.400 | 24.3 | 102.89 | 36206.378717 | 57227.840030 |
| 16330 | 48.7 | 68.927 | 8.511 | 7.816 | 14.1 | 142.317 | 9.791 | 26.3 | 109.09 | 11463.527923 | 34696.482826 |
| 53159 | 49.3 | 6.800 | -3.158 | 115.032 | 7.8 | 143.765 | 6.735 | 27.5 | 115.09 | 25801.322442 | 46554.912277 |

In [5]: `full_analysis(df_numeric)`

Correlation with SOG:



ANOVA:

| | sum_sq | df | F | PR(>F) |
|--------------------|--------------|---------|--------------|---------------|
| TWA_Abs | 50237.647144 | 1.0 | 69919.907652 | 0.000000e+00 |
| TWS | 1296.353938 | 1.0 | 1804.243486 | 0.000000e+00 |
| Side_lines | 711.900262 | 1.0 | 990.810745 | 1.504496e-212 |
| F_front | 705.964962 | 1.0 | 982.550094 | 7.807193e-211 |
| Trim | 431.924574 | 1.0 | 601.145317 | 6.488255e-131 |
| boat_weight | 275.959843 | 1.0 | 384.076243 | 9.205647e-85 |
| Line_C | 153.576337 | 1.0 | 213.744949 | 3.609029e-48 |
| Heel_Abs | 31.854273 | 1.0 | 44.334238 | 2.836645e-11 |
| ROT | 27.570528 | 1.0 | 38.372195 | 5.954434e-10 |
| F_back | 8.513609 | 1.0 | 11.849097 | 5.780156e-04 |
| Residual | 15106.520684 | 21025.0 | NaN | NaN |

Polynomial fit:
 $R^2: 0.876$
Used samples: 21036

Formula :
 $SOG \approx 24.463 + 0.049 * Heel_Abs + 0.336 * Line_C + 0.036 * ROT + 0.738 * Side_lines + -0.153 * Trim + 2.482 * TWA_Abs + 0.356 * TWS + -0.163 * boat_weight + 0.277 * F_front + 0.034 * F_back$

| | feature | coefficient |
|---|-------------|-------------|
| 5 | TWA_Abs | 2.482302 |
| 3 | Side_lines | 0.738299 |
| 6 | TWS | 0.355919 |
| 1 | Line_C | 0.335935 |
| 8 | F_front | 0.277414 |
| 7 | boat_weight | -0.162902 |
| 4 | Trim | -0.152871 |
| 0 | Heel_Abs | 0.048917 |
| 2 | ROT | 0.036293 |
| 9 | F_back | 0.034038 |

II Upwind:

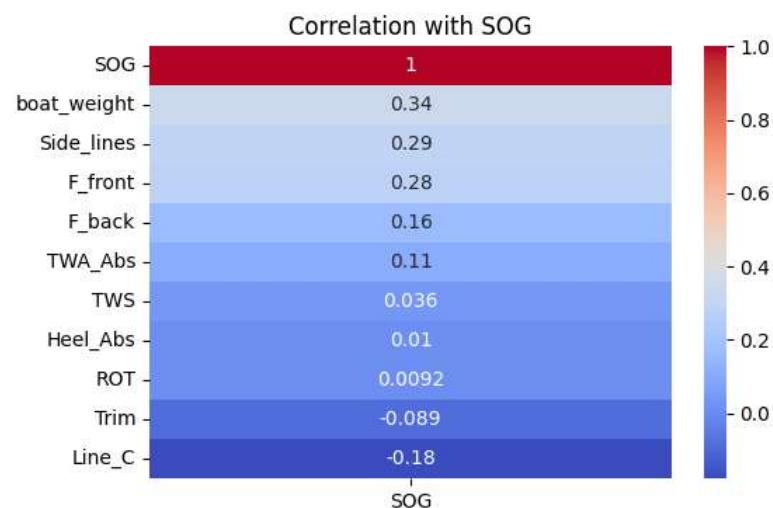
II.1. All upwind data

```
In [6]: upwind_data = df[df['TWA'] >= 0]
df_numeric_upwind = upwind_data.select_dtypes(include=["float64", "int64"]).copy()
df_numeric_upwind.drop(columns=[c for c in drop_cols if c in df_numeric_upwind.columns], inplace=True)
df_numeric_upwind.dropna(subset=["SOG"], inplace=True)
print(f"Number of rows after filtering: {len(df_numeric_upwind)}")
```

Number of rows after filtering: 18876

```
In [7]: full_analysis(df_numeric_upwind)
```

Correlation with SOG:



ANOVA:

| | sum_sq | df | F | PR(>F) |
|--------------------|-------------|---------|-------------|---------------|
| F_front | 745.352709 | 1.0 | 1412.437003 | 8.775914e-294 |
| TWS | 275.919193 | 1.0 | 522.864509 | 1.582006e-113 |
| Trim | 161.972730 | 1.0 | 306.936938 | 5.933708e-68 |
| Side_lines | 101.922591 | 1.0 | 193.142438 | 1.327238e-43 |
| TWA_Abs | 64.713004 | 1.0 | 122.630589 | 2.238728e-28 |
| F_back | 36.586519 | 1.0 | 69.331140 | 9.137337e-17 |
| Heel_Abs | 23.480768 | 1.0 | 44.495855 | 2.650378e-11 |
| ROT | 11.204361 | 1.0 | 21.232167 | 4.106812e-06 |
| Line_C | 1.860791 | 1.0 | 3.526183 | 6.042871e-02 |
| boat_weight | 0.355136 | 1.0 | 0.672980 | 4.120292e-01 |
| Residual | 6977.340227 | 13222.0 | NaN | NaN |

Polynomial fit:
 $R^2: 0.321$
Used samples: 13233

Formula :
 $SOG \approx 22.761 +$
 $-0.045 * Heel_Abs +$
 $0.050 * Line_C +$
 $0.030 * ROT +$
 $0.385 * Side_lines +$
 $-0.112 * Trim +$
 $0.082 * TWA_Abs +$
 $0.218 * TWS +$
 $0.007 * boat_weight +$
 $0.372 * F_front +$
 $0.093 * F_back$

| | feature | coefficient |
|----------|-------------|-------------|
| 3 | Side_lines | 0.384715 |
| 8 | F_front | 0.371517 |
| 6 | TWS | 0.217549 |
| 4 | Trim | -0.111691 |
| 9 | F_back | 0.093087 |
| 5 | TWA_Abs | 0.082455 |
| 1 | Line_C | 0.050125 |
| 0 | Heel_Abs | -0.045306 |
| 2 | ROT | 0.029856 |
| 7 | boat_weight | 0.007225 |

II.2. Upwind: Gian vs Karl

II.2.1. Upwind: Gian

```
In [8]: gian_data_upwind = upwind_data[
    (upwind_data['boat_name'] == "Gian Stragiotti") |
    ((upwind_data['boat_name'] == "SenseBoard") & (upwind_data['opponent_name'] == "Karl Maeder"))
]
df_numeric_gian_upwind = gian_data_upwind.select_dtypes(include=["float64", "int64"]).copy()
df_numeric_gian_upwind.drop(columns=[c for c in drop_cols if c in df_numeric_gian_upwind.columns], inplace=True)
df_numeric_gian_upwind.dropna(subset=["SOG"], inplace=True)
print(f"Number of rows after filtering: {len(df_numeric_gian_upwind)})")
```

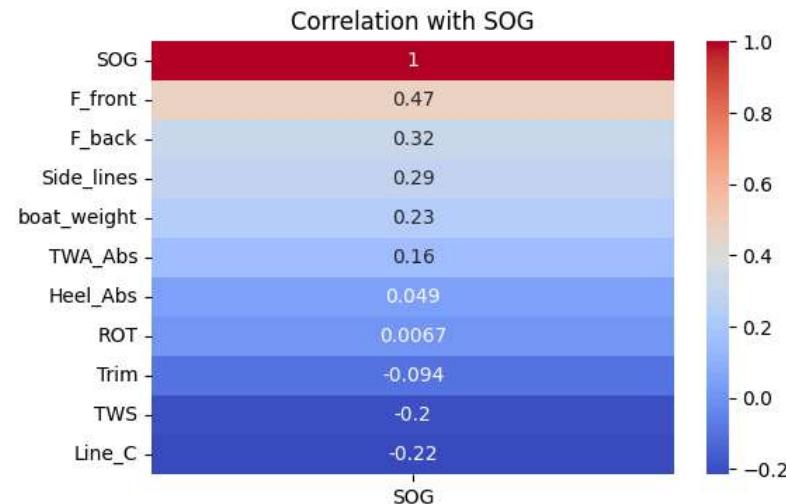
Number of rows after filtering: 13275

```
In [9]: print(gian_data_upwind["boat_weight"].mean)
```

```
<bound method Series.mean of 11633    109.09
11636    109.09
11637    109.09
11640    109.09
11641    109.09
...
80121    109.09
80124    109.09
80126    109.09
80127    109.09
80129    109.09
Name: boat_weight, Length: 13275, dtype: float64>
```

```
In [10]: full_analysis(df_numeric_gian_upwind)
```

Correlation with SOG:



ANOVA:

| | sum_sq | df | F | PR(>F) |
|--------------------|---------------|-----------|-------------|------------------|
| F_front | 745.352709 | 1.0 | 1412.437003 | 8.775914e-294 |
| TWS | 275.919193 | 1.0 | 522.864509 | 1.582006e-113 |
| Trim | 161.972730 | 1.0 | 306.936938 | 5.933708e-68 |
| Side_lines | 101.922591 | 1.0 | 193.142438 | 1.327238e-43 |
| TWA_Abs | 64.713004 | 1.0 | 122.630589 | 2.238728e-28 |
| F_back | 36.586519 | 1.0 | 69.331140 | 9.137337e-17 |
| Heel_Abs | 23.480768 | 1.0 | 44.495855 | 2.650378e-11 |
| ROT | 11.204361 | 1.0 | 21.232167 | 4.106812e-06 |
| Line_C | 1.860791 | 1.0 | 3.526183 | 6.042871e-02 |
| boat_weight | 0.355136 | 1.0 | 0.672980 | 4.120292e-01 |
| Residual | 6977.340227 | 13222.0 | NaN | NaN |

Polynomial fit:

R²: 0.321

Used samples: 13233

```
Formula :
SOG ~ 22.761 +
-0.045 * Heel_Abs +
0.050 * Line_C +
0.030 * ROT +
0.385 * Side_lines +
-0.112 * Trim +
0.082 * TWA_Abs +
0.218 * TWS +
0.007 * boat_weight +
0.372 * F_front +
0.093 * F_back
```

| | feature | coefficient |
|----------|----------------|--------------------|
| 3 | Side_lines | 0.384715 |
| 8 | F_front | 0.371517 |
| 6 | TWS | 0.217549 |
| 4 | Trim | -0.111691 |
| 9 | F_back | 0.093087 |
| 5 | TWA_Abs | 0.082455 |
| 1 | Line_C | 0.050125 |
| 0 | Heel_Abs | -0.045306 |
| 2 | ROT | 0.029856 |
| 7 | boat_weight | 0.007225 |

II.2.2. Upwind: Karl

```
In [11]: karl_data_upwind = upwind_data[((upwind_data['boat_name'] == "SenseBoard") & (upwind_data['opponent_name'] == "Gian Stragiotti"))
]
df_numeric_karl_upwind = karl_data_upwind.select_dtypes(include=["float64", "int64"]).copy()
df_numeric_karl_upwind.drop(columns=[c for c in drop_cols if c in df_numeric_karl_upwind.columns], inplace=True)
df_numeric_karl_upwind.dropna(subset=["SOG"], inplace=True)
print(f"Number of rows after filtering: {len(df_numeric_karl_upwind)}")
print(df_numeric_karl_upwind.shape)
print(df_numeric_karl_upwind.columns)
```

Number of rows after filtering: 5601
(5601, 11)
Index(['Heel_Abs', 'Line_C', 'ROT', 'Side_lines', 'Trim', 'TWA_Abs', 'TWS',
'SOG', 'boat_weight', 'F_front', 'F_back'],
dtype='object')

```
In [12]: print(karl_data_upwind["boat_weight"].mean)
```

<bound method Series.mean of 25553 102.89
25555 102.89
25557 102.89
25560 102.89
25562 102.89
...
68904 102.89
68907 102.89
68908 102.89
68911 102.89
68912 102.89
Name: boat_weight, Length: 5601, dtype: float64>

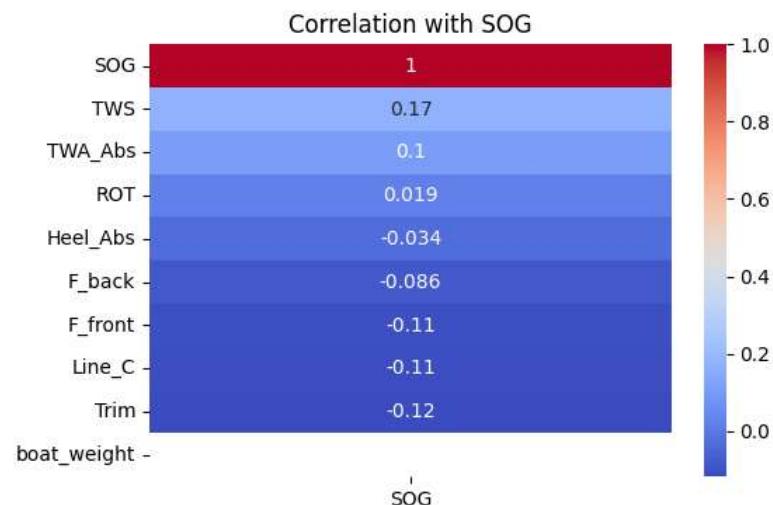
```
In [13]: df_numeric_karl_upwind.head(10)
```

| | Heel_Abs | Line_C | ROT | Side_lines | Trim | TWA_Abs | TWS | SOG | boat_weight | F_front | F_back |
|--------------|----------|--------|---------|------------|------|---------|-------|------|-------------|--------------|--------------|
| 25553 | 59.5 | 93.3 | 11.000 | NaN | 7.5 | 41.174 | 9.352 | 21.4 | 102.89 | 27567.598622 | 24661.188272 |
| 25555 | 57.4 | 88.1 | -14.141 | NaN | 12.3 | 42.564 | 9.372 | 21.4 | 102.89 | 30282.257552 | 24532.446152 |
| 25557 | 59.8 | 98.4 | 1.020 | NaN | 8.7 | 42.454 | 9.391 | 21.3 | 102.89 | 30629.872010 | 26538.607253 |
| 25560 | 57.8 | 108.1 | 15.000 | NaN | 1.5 | 40.944 | 9.411 | 21.4 | 102.89 | 29534.445709 | 25671.679650 |
| 25562 | 53.0 | 111.4 | -7.619 | NaN | 4.6 | 41.734 | 9.432 | 21.4 | 102.89 | 31969.060202 | 27080.840797 |
| 25563 | 54.8 | 111.3 | -10.309 | NaN | 7.0 | 42.724 | 9.452 | 21.4 | 102.89 | 32674.719191 | 26459.098753 |
| 25565 | 54.6 | 110.7 | 11.224 | NaN | 7.0 | 41.614 | 9.471 | 21.3 | 102.89 | 34743.378192 | 26875.356725 |
| 25567 | 47.0 | 109.9 | -19.802 | NaN | 6.2 | 43.604 | 9.492 | 21.4 | 102.89 | 34743.378192 | 26875.356725 |
| 25569 | 41.2 | 107.5 | 7.000 | NaN | 4.4 | 42.894 | 9.506 | 21.3 | 102.89 | 30655.266332 | 26253.937186 |
| 25571 | 51.7 | 104.6 | 1.000 | NaN | 8.5 | 42.784 | 9.516 | 21.4 | 102.89 | 30655.266332 | 26253.937186 |

```
In [14]: full_analysis(df_numeric_karl_upwind)
```

Dropped column 'Side_lines' (all values were NaN).

Correlation with SOG:



ANOVA:

| | sum_sq | df | F | PR(>F) |
|--------------------|-------------|--------|-------------|---------------|
| boat_weight | 4598.244751 | 1.0 | 9685.361760 | 0.000000e+00 |
| TWA_Abs | 246.517194 | 1.0 | 519.243393 | 5.706783e-110 |
| TWS | 195.231855 | 1.0 | 411.220205 | 2.842816e-88 |
| Trim | 93.653566 | 1.0 | 197.264112 | 4.622077e-44 |
| F_back | 15.631601 | 1.0 | 32.925110 | 1.008229e-08 |
| ROT | 12.608568 | 1.0 | 26.557643 | 2.646028e-07 |
| Heel_Abs | 9.943821 | 1.0 | 20.944841 | 4.829555e-06 |
| Line_C | 5.844075 | 1.0 | 12.309476 | 4.542211e-04 |
| F_front | 2.640326 | 1.0 | 5.561364 | 1.839543e-02 |
| Residual | 2646.799895 | 5575.0 | NaN | NaN |

```
Polynomial fit:
R²: 0.135
Used samples: 5584

Formula :
SOG ≈ 22.227 +
0.046 * Heel_Abs +
-0.045 * Line_C +
0.049 * ROT +
-0.136 * Trim +
0.274 * TWA_Abs +
0.336 * TWS +
0.000 * boat_weight +
-0.029 * F_front +
0.086 * F_back
```

| | feature | coefficient |
|---|-------------|---------------|
| 5 | TWS | 3.358031e-01 |
| 4 | TWA_Abs | 2.736814e-01 |
| 3 | Trim | -1.357438e-01 |
| 8 | F_back | 8.641826e-02 |
| 2 | ROT | 4.868985e-02 |
| 0 | Heel_Abs | 4.570711e-02 |
| 1 | Line_C | -4.452198e-02 |
| 7 | F_front | -2.866813e-02 |
| 6 | boat_weight | 5.551115e-17 |

II.2.3. Upwind: Karl vs Gian t-test

```
In [15]: t_test(df_numeric_gian_upwind, df_numeric_karl_upwind)

T-statistic: 39.897, p-value: 0.0000000000000000
The difference is statistically significant, keeping data split.
```

II.3. Upwind: Master vs Slave

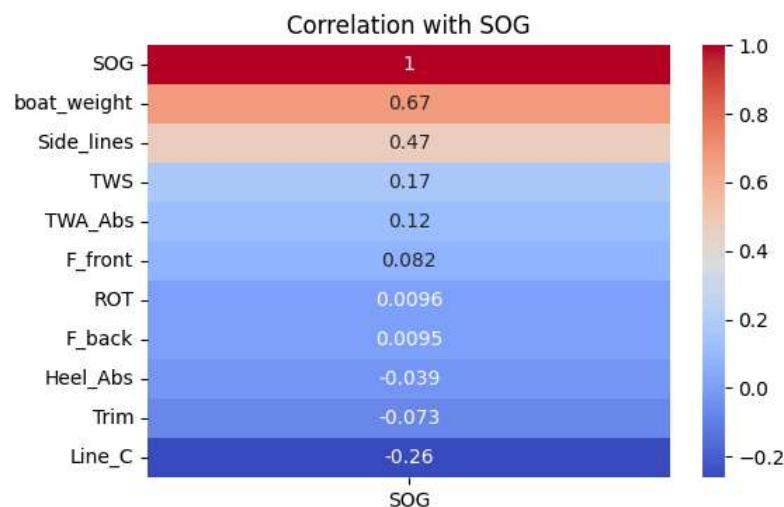
II.3.1. Master

```
In [16]: master_data_upwind = upwind_data[upwind_data['boat_role'] == "master"]
df_numeric_master_upwind = master_data_upwind.select_dtypes(include=["float64", "int64"]).copy()
df_numeric_master_upwind.drop(columns=[c for c in drop_cols if c in df_numeric_master_upwind.columns], inplace=True)
df_numeric_master_upwind.dropna(subset=["SOG"], inplace=True)
print(f"Number of rows after filtering: {len(df_numeric_master_upwind)}")

Number of rows after filtering: 9638
```

```
In [17]: full_analysis(df_numeric_master_upwind)
```

Correlation with SOG:



ANOVA:

| | sum_sq | df | F | PR(>F) |
|--------------------|-------------|--------|------------|---------------|
| F_front | 168.253949 | 1.0 | 480.684713 | 1.004899e-102 |
| Side_lines | 70.414496 | 1.0 | 201.167177 | 5.723545e-45 |
| Trim | 50.526690 | 1.0 | 144.349702 | 6.826664e-33 |
| TWS | 18.385077 | 1.0 | 52.524329 | 4.760053e-13 |
| Heel_Abs | 8.065963 | 1.0 | 23.043649 | 1.620350e-06 |
| boat_weight | 7.376923 | 1.0 | 21.075133 | 4.502315e-06 |
| Line_C | 4.703126 | 1.0 | 13.436360 | 2.488331e-04 |
| F_back | 1.449286 | 1.0 | 4.140464 | 4.191165e-02 |
| ROT | 0.959391 | 1.0 | 2.740884 | 9.786063e-02 |
| TWA_Abs | 0.929970 | 1.0 | 2.656830 | 1.031562e-01 |
| Residual | 2196.086636 | 6274.0 | Nan | Nan |

Polynomial fit:

R²: 0.358

Used samples: 6285

```
Formula :
SOG ≈ 23.170 +
-0.039 * Heel_Abs +
0.110 * Line_C +
0.013 * ROT +
0.445 * Side_lines +
-0.091 * Trim +
-0.017 * TWA_Abs +
0.091 * TWS +
0.059 * boat_weight +
0.270 * F_front +
0.029 * F_back
```

| | feature | coefficient |
|----------|-------------|-------------|
| 3 | Side_lines | 0.445396 |
| 8 | F_front | 0.270209 |
| 1 | Line_C | 0.110405 |
| 6 | TWS | 0.091217 |
| 4 | Trim | -0.090709 |
| 7 | boat_weight | 0.058505 |
| 0 | Heel_Abs | -0.038598 |
| 9 | F_back | 0.028592 |
| 5 | TWA_Abs | -0.017015 |
| 2 | ROT | 0.012810 |

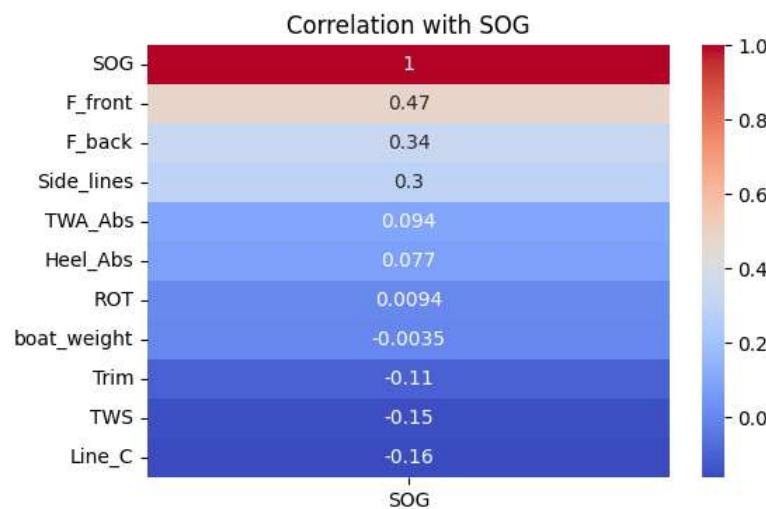
II.3.2 Slave

```
In [18]: slave_data_upwind = upwind_data[upwind_data['boat_role'] == "slave"]
df_numeric_slave_upwind = slave_data_upwind.select_dtypes(include=["float64", "int64"]).copy()
df_numeric_slave_upwind.drop(columns=[c for c in drop_cols if c in df_numeric_slave_upwind.columns], inplace=True)
df_numeric_slave_upwind.dropna(subset=["SOG"], inplace=True)
print(f"Number of rows after filtering: {len(df_numeric_slave_upwind)})")
```

Number of rows after filtering: 9238

```
In [19]: full_analysis(df_numeric_slave_upwind)
```

Correlation with SOG:



ANOVA:

| | sum_sq | df | F | PR(>F) |
|--------------------|-------------|--------|------------|---------------|
| F_front | 258.589307 | 1.0 | 616.290151 | 2.058431e-130 |
| F_back | 85.092555 | 1.0 | 202.799196 | 2.221636e-45 |
| Trim | 63.317539 | 1.0 | 150.903284 | 2.497524e-34 |
| Side_lines | 48.725358 | 1.0 | 116.126064 | 7.272527e-27 |
| TWS | 42.270916 | 1.0 | 100.743336 | 1.515060e-23 |
| boat_weight | 33.370761 | 1.0 | 79.531794 | 5.983729e-19 |
| Heel_Abs | 22.101350 | 1.0 | 52.673657 | 4.366700e-13 |
| TWA_Abs | 10.436726 | 1.0 | 24.873618 | 6.270176e-07 |
| Line_C | 0.323681 | 1.0 | 0.771422 | 3.798080e-01 |
| ROT | 0.067776 | 1.0 | 0.161529 | 6.877648e-01 |
| Residual | 2910.697207 | 6937.0 | Nan | Nan |

Polynomial fit:
 $R^2: 0.401$
Used samples: 6948

Formula :
 $SOG \approx 22.390 +$
 $-0.061 * Heel_Abs +$
 $0.030 * Line_C +$
 $0.003 * ROT +$
 $0.381 * Side_lines +$
 $-0.097 * Trim +$
 $-0.048 * TWA_Abs +$
 $0.134 * TWS +$
 $-0.091 * boat_weight +$
 $0.315 * F_front +$
 $0.200 * F_back$

| | feature | coefficient |
|----------|-------------|-------------|
| 3 | Side_lines | 0.380791 |
| 8 | F_front | 0.315048 |
| 9 | F_back | 0.200336 |
| 6 | TWS | 0.133552 |
| 4 | Trim | -0.096585 |
| 7 | boat_weight | -0.090734 |
| 0 | Heel_Abs | -0.060942 |
| 5 | TWA_Abs | -0.047994 |
| 1 | Line_C | 0.030051 |
| 2 | ROT | 0.003210 |

III.3.3. Upwind: Master vs Slave t-test

```
In [20]: t_test(df_numeric_master_upwind, df_numeric_slave_upwind)

T-statistic: 25.880, p-value: 0.0000000000000000
The difference is statistically significant, keeping data split.
```

III Downwind

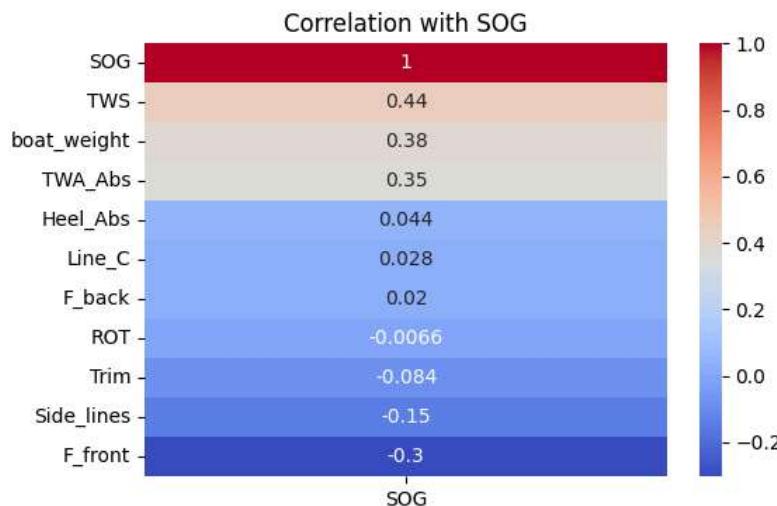
III.1. All downwind data

```
In [21]: downwind_data = df[df['TWA'] < 0]
df_numeric_downwind = downwind_data.select_dtypes(include=["float64", "int64"]).copy()
df_numeric_downwind.drop(columns=[c for c in drop_cols if c in df_numeric_downwind.columns], inplace=True)
df_numeric_downwind.dropna(subset=['SOG'], inplace=True)
print(f"Number of rows after filtering: {len(df_numeric_downwind)}")
```

Number of rows after filtering: 11670

```
In [22]: full_analysis(df_numeric_downwind)
```

Correlation with SOG:



ANOVA:

| | sum_sq | df | F | PR(>F) |
|-------------|-------------|--------|-------------|---------------|
| TWS | 962.827764 | 1.0 | 1191.018987 | 5.314048e-243 |
| Side_lines | 379.737135 | 1.0 | 469.735247 | 3.407176e-101 |
| boat_weight | 209.666280 | 1.0 | 259.357416 | 1.989188e-57 |
| Line_C | 144.803508 | 1.0 | 179.122096 | 2.100678e-40 |
| Trim | 133.760353 | 1.0 | 165.461702 | 1.742471e-37 |
| Heel_Abs | 129.867362 | 1.0 | 160.646067 | 1.869130e-36 |
| TWA_Abs | 74.036677 | 1.0 | 91.583450 | 1.405424e-21 |
| F_back | 36.755196 | 1.0 | 45.466217 | 1.663794e-11 |
| F_front | 33.688069 | 1.0 | 41.672178 | 1.143975e-10 |
| ROT | 1.926524 | 1.0 | 2.383113 | 1.226936e-01 |
| Residual | 6299.105237 | 7792.0 | NaN | NaN |

Polynomial fit:
 $R^2: 0.262$
Used samples: 7803

Formula :
 $SOG \approx 27.350 + 0.174 * Heel_Abs + 0.503 * Line_C + -0.016 * ROT + 0.828 * Side_lines + -0.158 * Trim + 0.124 * TWA_Abs + 0.498 * TWS + -0.264 * boat_weight + -0.100 * F_front + 0.112 * F_back$

| | feature | coefficient |
|----------|-------------|-------------|
| 3 | Side_lines | 0.827594 |
| 1 | Line_C | 0.503088 |
| 6 | TWS | 0.497873 |
| 7 | boat_weight | -0.264347 |
| 0 | Heel_Abs | 0.173604 |
| 4 | Trim | -0.158352 |
| 5 | TWA_Abs | 0.124349 |
| 9 | F_back | 0.111924 |
| 8 | F_front | -0.100203 |
| 2 | ROT | -0.016486 |

III.2. Downwind: Gian vs Karl

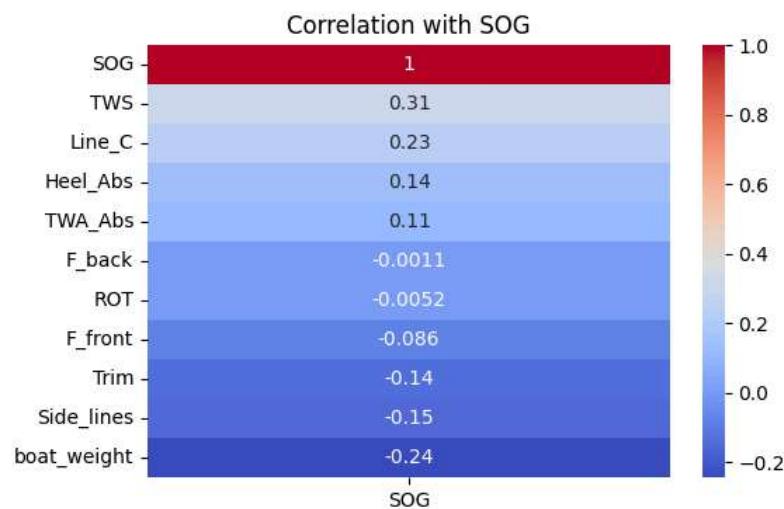
III.2.1. Downwind: Gian

```
In [23]: gian_data_downwind = downwind_data[
    (downwind_data['boat_name'] == "Gian Stragiotti") |
    ((downwind_data['boat_name'] == "SenseBoard") & (downwind_data['opponent_name'] == "Karl Maeder"))
]
df_numeric_gian_downwind = gian_data_downwind.select_dtypes(include=["float64", "int64"]).copy()
df_numeric_gian_downwind.drop(columns=[c for c in drop_cols if c in df_numeric_gian_downwind.columns], inplace=True)
df_numeric_gian_downwind.dropna(subset=["SOG"], inplace=True)
print(f"Number of rows after filtering: {len(df_numeric_gian_downwind)}")
```

Number of rows after filtering: 7826

```
In [24]: full_analysis(df_numeric_gian_downwind)
```

Correlation with SOG:



ANOVA:

| | sum_sq | df | F | PR(>F) |
|--------------------|-------------|--------|-------------|---------------|
| TWS | 962.827764 | 1.0 | 1191.018987 | 5.314048e-243 |
| Side_lines | 379.737135 | 1.0 | 469.735247 | 3.407176e-101 |
| boat_weight | 209.666280 | 1.0 | 259.357416 | 1.989188e-57 |
| Line_C | 144.803508 | 1.0 | 179.122096 | 2.100678e-40 |
| Trim | 133.760353 | 1.0 | 165.461702 | 1.742471e-37 |
| Heel_Abs | 129.867362 | 1.0 | 160.646067 | 1.869130e-36 |
| TWA_Abs | 74.036677 | 1.0 | 91.583450 | 1.405424e-21 |
| F_back | 36.755196 | 1.0 | 45.466217 | 1.663794e-11 |
| F_front | 33.688069 | 1.0 | 41.672178 | 1.143975e-10 |
| ROT | 1.926524 | 1.0 | 2.383113 | 1.226936e-01 |
| Residual | 6299.105237 | 7792.0 | NaN | NaN |

Polynomial fit:
 $R^2: 0.262$
Used samples: 7803

Formula :
 $SOG \approx 27.350 + 0.174 * Heel_Abs + 0.503 * Line_C + -0.016 * ROT + 0.828 * Side_lines + -0.158 * Trim + 0.124 * TWA_Abs + 0.498 * TWS + -0.264 * boat_weight + -0.100 * F_front + 0.112 * F_back$

| | feature | coefficient |
|----------|-------------|-------------|
| 3 | Side_lines | 0.827594 |
| 1 | Line_C | 0.503088 |
| 6 | TWS | 0.497873 |
| 7 | boat_weight | -0.264347 |
| 0 | Heel_Abs | 0.173604 |
| 4 | Trim | -0.158352 |
| 5 | TWA_Abs | 0.124349 |
| 9 | F_back | 0.111924 |
| 8 | F_front | -0.100203 |
| 2 | ROT | -0.016486 |

III.2.2. Downwind: Karl

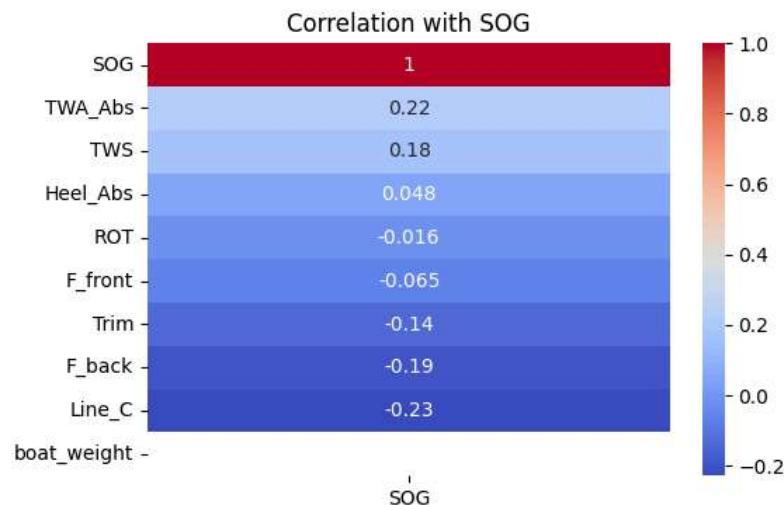
```
In [25]: karl_data_downwind = downwind_data[
    (downwind_data['boat_name'] == "Karl Maeder") |
    ((downwind_data['boat_name'] == "SenseBoard") & (downwind_data['opponent_name'] == "Gian Stragiotti"))
]
df_numeric_karl_downwind = karl_data_downwind.select_dtypes(include=["float64", "int64"]).copy()
df_numeric_karl_downwind.drop(columns=[c for c in drop_cols if c in df_numeric_karl_downwind.columns], inplace=True)
df_numeric_karl_downwind.dropna(subset=["SOG"], inplace=True)
print(f"Number of rows after filtering: {len(df_numeric_karl_downwind)}")
```

Number of rows after filtering: 3844

```
In [26]: full_analysis(df_numeric_karl_downwind)
```

Dropped column 'Side_lines' (all values were NaN).

Correlation with SOG:



ANOVA:

| | sum_sq | df | F | PR(>F) |
|--------------------|-------------|--------|-------------|--------------|
| boat_weight | 1978.470978 | 1.0 | 3024.729178 | 0.000000e+00 |
| Trim | 292.014819 | 1.0 | 446.438565 | 8.025132e-94 |
| Heel_Abs | 177.459225 | 1.0 | 271.303498 | 6.013933e-59 |
| Line_C | 161.363777 | 1.0 | 246.696429 | 6.355927e-54 |
| F_front | 133.389888 | 1.0 | 203.929342 | 4.098325e-45 |
| TWA_Abs | 55.567557 | 1.0 | 84.952882 | 4.914748e-20 |
| F_back | 16.445748 | 1.0 | 25.142614 | 5.565583e-07 |
| TWS | 11.771670 | 1.0 | 17.996784 | 2.265131e-05 |
| ROT | 6.755128 | 1.0 | 10.327385 | 1.321582e-03 |
| Residual | 2504.543358 | 3829.0 | Nan | Nan |

Polynomial fit:

R²: 0.230

Used samples: 3838

```

Formula :
SOG ≈ 25.975 +
0.258 * Heel_Abs +
-0.332 * Line_C +
-0.045 * ROT +
-0.325 * Trim +
0.204 * TWA_Abs +
0.090 * TWS +
0.000 * boat_weight +
0.462 * F_front +
-0.231 * F_back

```

| | feature | coefficient |
|---|-------------|---------------|
| 7 | F_front | 4.622215e-01 |
| 1 | Line_C | -3.324574e-01 |
| 3 | Trim | -3.247854e-01 |
| 0 | Heel_Abs | 2.577719e-01 |
| 8 | F_back | -2.305324e-01 |
| 4 | TWA_Abs | 2.043650e-01 |
| 5 | TWS | 9.002057e-02 |
| 2 | ROT | -4.452262e-02 |
| 6 | boat_weight | 3.330669e-16 |

III.2.3. Downwind: Karl vs Gian t-test

In [27]: `t_test(df_numeric_karl_downwind, df_numeric_gian_downwind)`

T-statistic: -69.413, p-value: 0.0000000000000000
The difference is statistically significant, keeping data split.

III.3. Downwind: Master vs Slave

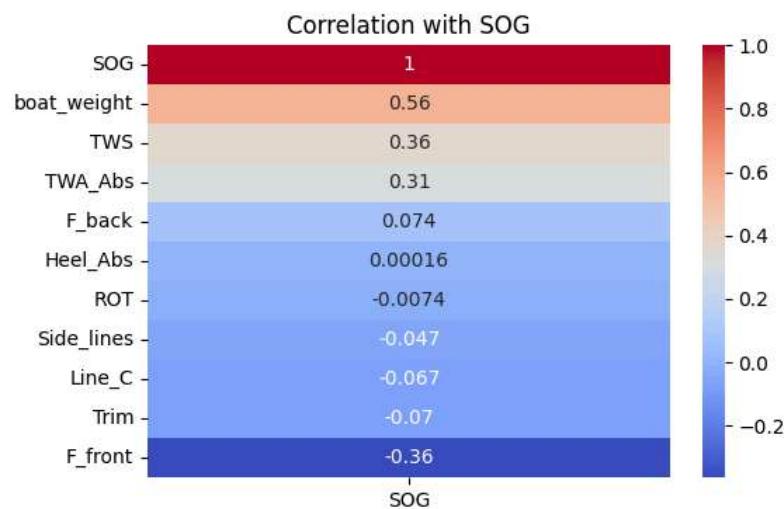
III.3.1 Downwind Master

In [28]: `master_data_downwind = downwind_data[downwind_data['boat_role'] == "master"]
df_numeric_master_downwind = master_data_downwind.select_dtypes(include=["float64", "int64"]).copy()
df_numeric_master_downwind.drop(columns=[c for c in drop_cols if c in df_numeric_master_downwind.columns], inplace=True)
df_numeric_master_downwind.dropna(subset=["SOG"], inplace=True)
print(f"Number of rows after filtering: {len(df_numeric_master_downwind)}")`

Number of rows after filtering: 6136

In [29]: `full_analysis(df_numeric_master_downwind)`

Correlation with SOG:



ANOVA:

| | sum_sq | df | F | PR(>F) |
|--------------------|-------------|--------|------------|--------------|
| TWS | 182.630180 | 1.0 | 308.683119 | 1.885832e-66 |
| TWA_Abs | 74.777435 | 1.0 | 126.389471 | 7.328615e-29 |
| Trim | 53.496760 | 1.0 | 90.420690 | 3.343194e-21 |
| Side_lines | 50.591508 | 1.0 | 85.510207 | 3.779641e-20 |
| Heel_Abs | 44.092802 | 1.0 | 74.526037 | 8.729937e-18 |
| boat_weight | 19.757555 | 1.0 | 33.394392 | 8.140845e-09 |
| F_front | 18.181873 | 1.0 | 30.731160 | 3.169275e-08 |
| Line_C | 5.539048 | 1.0 | 9.362147 | 2.230889e-03 |
| F_back | 4.718138 | 1.0 | 7.974639 | 4.768922e-03 |
| ROT | 2.745746 | 1.0 | 4.640884 | 3.128283e-02 |
| Residual | 2204.461489 | 3726.0 | NaN | NaN |

Polynomial fit:
 $R^2: 0.183$
Used samples: 3737

Formula :
 $SOG \approx 27.537 + 0.143 * Heel_Abs + 0.146 * Line_C + -0.028 * ROT + 0.478 * Side_lines + -0.143 * Trim + 0.179 * TWA_Abs + 0.400 * TWS + -0.138 * boat_weight + -0.185 * F_front + 0.061 * F_back$

| | feature | coefficient |
|---|-------------|-------------|
| 3 | Side_lines | 0.478106 |
| 6 | TWS | 0.400064 |
| 5 | TWA_Abs | 0.179439 |
| 1 | Line_C | 0.145757 |
| 4 | Trim | -0.142889 |
| 0 | Heel_Abs | 0.142679 |
| 7 | boat_weight | -0.138464 |
| 8 | F_front | -0.104683 |
| 9 | F_back | 0.060866 |
| 2 | ROT | -0.028425 |

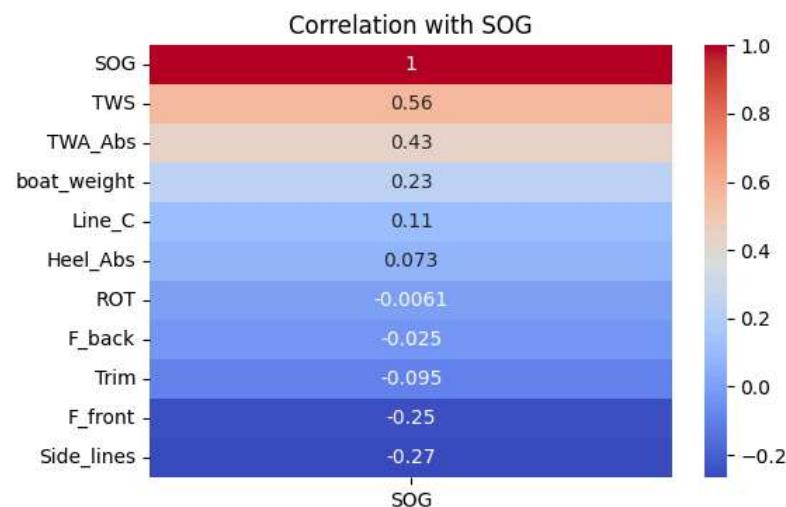
III.3.2 Downwind Slave

```
In [30]: slave_data_downwind = downwind_data[downwind_data['boat_role'] == "slave"]
df_numeric_slave_downwind = slave_data_downwind.select_dtypes(include=["float64", "int64"]).copy()
df_numeric_slave_downwind.drop(columns=[c for c in drop_cols if c in df_numeric_slave_downwind.columns], inplace=True)
df_numeric_slave_downwind.dropna(subset=['SOG'], inplace=True)
print(f"Number of rows after filtering: {len(df_numeric_slave_downwind)}")
```

Number of rows after filtering: 5534

```
In [31]: full_analysis(df_numeric_slave_downwind)
```

Correlation with SOG:



ANOVA:

| | sum_sq | df | F | PR(>F) |
|--------------------|-------------|--------|------------|--------------|
| TWS | 348.757145 | 1.0 | 385.778206 | 4.032141e-82 |
| Side_lines | 238.573968 | 1.0 | 263.898929 | 1.534198e-57 |
| Line_C | 163.060321 | 1.0 | 180.369403 | 2.885847e-40 |
| boat_weight | 67.939673 | 1.0 | 75.151565 | 6.205239e-18 |
| Trim | 55.352436 | 1.0 | 61.228174 | 6.437538e-15 |
| Heel_Abs | 46.431791 | 1.0 | 51.360591 | 9.088353e-13 |
| F_back | 30.010050 | 1.0 | 33.195658 | 8.952174e-09 |
| TWA_Abs | 16.079200 | 1.0 | 17.786029 | 2.525872e-05 |
| F_front | 7.832009 | 1.0 | 8.663388 | 3.265090e-03 |
| ROT | 0.131462 | 1.0 | 0.145417 | 7.029743e-01 |
| Residual | 3665.863451 | 4055.0 | NaN | NaN |

Polynomial fit:

R²: 0.344

Used samples: 4066

Formula :

$$\text{SOG} \approx 27.178 + 0.147 * \text{Heel_Abs} + 0.739 * \text{Line_C} + -0.006 * \text{ROT} + 0.871 * \text{Side_lines} + -0.145 * \text{Trim} + 0.083 * \text{TWA_Abs} + 0.464 * \text{TWS} + -0.246 * \text{boat_weight} + -0.072 * \text{F_front} + 0.139 * \text{F_back}$$

| | feature | coefficient |
|----------|-------------|-------------|
| 3 | Side_lines | 0.870608 |
| 1 | Line_C | 0.739499 |
| 6 | TWS | 0.463966 |
| 7 | boat_weight | -0.246087 |
| 0 | Heel_Abs | 0.147356 |
| 4 | Trim | -0.145263 |
| 9 | F_back | 0.139004 |
| 5 | TWA_Abs | 0.083093 |
| 8 | F_front | -0.071689 |
| 2 | ROT | -0.005988 |

III.3.3. Downwind: Master vs Slave t-test

In [32]: `t_test(df_numeric_master_downwind,df_numeric_slave_downwind)`

T-statistic: 9.267, p-value: 0.000000000000000

The difference is statistically significant, keeping data split.